



5 Object detection

2024年9月26日 16:13

1 Task:



2.



Bounding box: visible part

Amodal detection: entire (occluded parts)

3 Intersection over union (IoU)

or Jaccard similarity / Jaccard index

$$= \frac{S_i}{S_u} \quad \text{intersection over union}$$

Pred v.s. ground-truth box (GT):


$IoU > 0.5$ decent

> 0.7 pretty good


> 0.9 almost perfect

4 Models

1° To detect a single obj.



2° To det multiple obj




↓ # boxes

⇒ Region Proposals

- Find a small set of boxes that are likely to cover all objects
- Often based on heuristics: e.g. look for "blob-like" image regions
- Relatively fast to run; e.g. Selective Search gives 2000 region proposals in a few seconds on CPU



3° R-CNN (Region-CNN)



R-CNN: Box Regression



Consider a region proposal with center (p_x, p_y) , with p_w, p_h .

Model predicts $\hat{p}_x, \hat{p}_y, \hat{p}_w, \hat{p}_h$.

The output box is defined by:

$\hat{p}_x = p_x + \hat{p}_x$

$\hat{p}_y = p_y + \hat{p}_y$

$\hat{p}_w = p_w \exp(\hat{p}_w)$

$\hat{p}_h = p_h \exp(\hat{p}_h)$

When transform is 0, output = proposal

L2 regularization encourages leaving proposal unchanged

we can also use the following formula:

$\hat{p}_x = p_x + \hat{p}_x$

$\hat{p}_y = p_y + \hat{p}_y$

$\hat{p}_w = p_w \exp(\hat{p}_w)$

$\hat{p}_h = p_h \exp(\hat{p}_h)$

How can we compare our prediction to the ground-truth box?


Intersection over Union (IoU)

(Also called "Jaccard similarity" or "Jaccard index"):

Area of Intersection

Area of Union

$IoU > 0.5$ is "decent"



Pb1

⇒ Non-Max Suppression solve overlapping boxes

消除IoU 太高, scoring 低 ~ 那个直到完全 > threshold.

⊖ Obj overlapping

1. Select max highest scoring box
2. Find other boxes with higher IoU than threshold
3. If any, delete them, GO TO 1



Pb2: L2 reg encourages leaving proposal unchanged
sol: mean Avg. Precision (mAP)

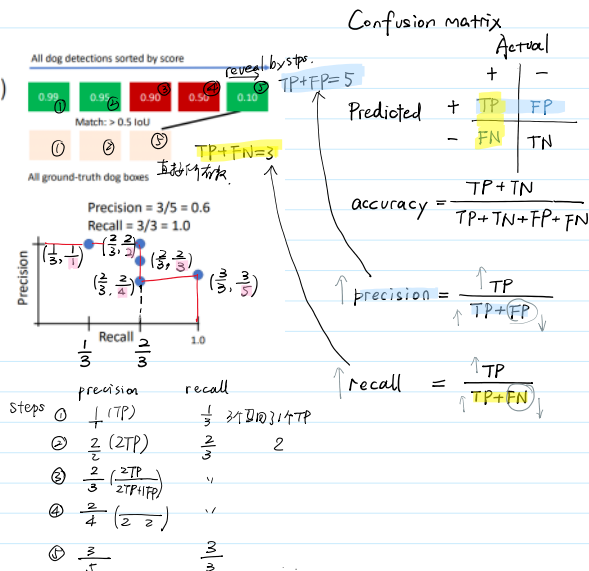
Evaluating Object Detectors: Mean Average Precision (mAP)

- Run object detector on all test images (with NMS)
- For each category, compute Average Precision (AP) = area under Precision vs Recall Curve
 - For each detection (highest score to lowest score)
 - If it matches some GT box with IoU > 0.5 , mark it as positive and eliminate the GT
 - Otherwise mark it as negative
 - Plot a point on PR Curve

$$AP = \int_0^1 p(r) dr$$

$$= \frac{2}{3} \times 1 + \frac{0.5 \times 0.4}{2} \times \frac{1}{3}$$

$$= \frac{2}{3} + \frac{3}{60} = \frac{40+9}{60} = \frac{49}{60} = 0.81$$



- data balanced
- easy to exp
- 每类 label 一样重要


查准率: 在预测为T中, 有多少真T. 求精.

① 冤假错案成本高: 垃圾邮件.

查全: 实际为T的有多少被预测到.

① 漏网之鱼成本高: 地震预警.

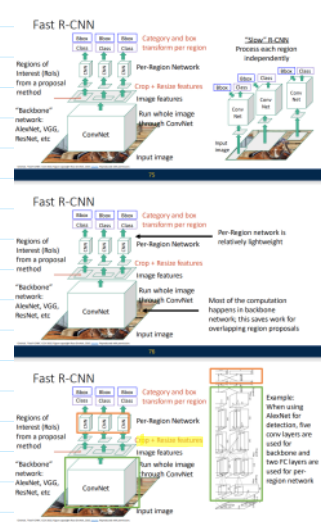
Pb3 Slow.



下一个字母+1: if TP, 分子+1.

字母-1为 TP+FN: if TP, 分子+1

4* Fast R-CNN



Crop features

1° RoI Pool

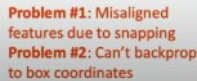
different resolution here

Cropping Features: RoI Pool

Divide into 2x2

RoI Pool

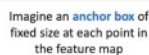
1° Rol Pool




bilinear interpolation



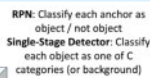
can be backpropagate to box coordinates!



+




Problem: Anchor box may have the wrong size / shape
Solution: Use K different anchor boxes at each point!



a little better




- Two stage method (Faster R-CNN) get the best accuracy, but are slower
- Single-stage methods (SSD) are much faster, but don't perform as well
- Bigger backbones improve performance, but are slower




Object Detection Methods

Both of these rely on anchor boxes.
Can we do detection without anchors?

"Slow" R-CNN: Run CNN independently for each region



Fast R-CNN: Apply differentiable cropping to shared image features



Faster R-CNN: Compute proposals with CNN



Single-Stage: Fully convolutional detector



Yes cornerNet